

PFEAST: A High Performance Sparse Eigenvalue Solver Using Distributed-Memory Linear Solvers



James Kestyn, Vasileios Kalantzis, Eric Polizzi, and Yousef Saad



Tuesday, November 15th, 2016

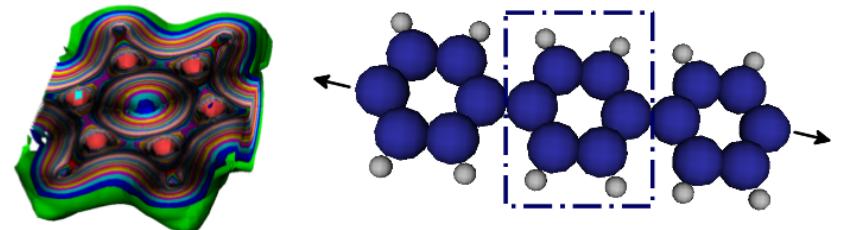
Eigenvalue Problems

$$AX = BX\Lambda$$

Ubiquitous in science and engineering with a diverse range of applications from Physics and Chemistry to Numerical Analysis and Data Analytics

- Principal Component Analysis (Statistics ...)
- Stability of Electrical Networks
- Vibrational Analysis (Structural Engineering)
- Dimensionality Reduction

Electronic Structure



- Electron energies/configuration within a material (i.e. bonding)
- Used to predict and verify material properties from first principles

Eigenvalue Algorithms/Solvers

Dense Problems

- Small matrices
- Compute the entire spectrum
- Efficient solvers in LAPACK

Sparse Problems

- Large matrices
- Interior Eigenproblem (subset of eigenpairs)
- Many different algorithms exist

Sparse Eigensolvers

Krylov Subspace (Arnoldi, Lanczos)

ARPACK, FITLAN and TRLAN

Jacobi-Davidson

PRIMME and JADAMILU

LOBPCG

BLOPEX

Other Libraries and Algorithms

SLEPc and ANASAZI

TraceMin

Contour Integration Techniques

zPARES and FEAST

Difficulty: must compute slices of spectrum independently



The FEAST Eigensolver

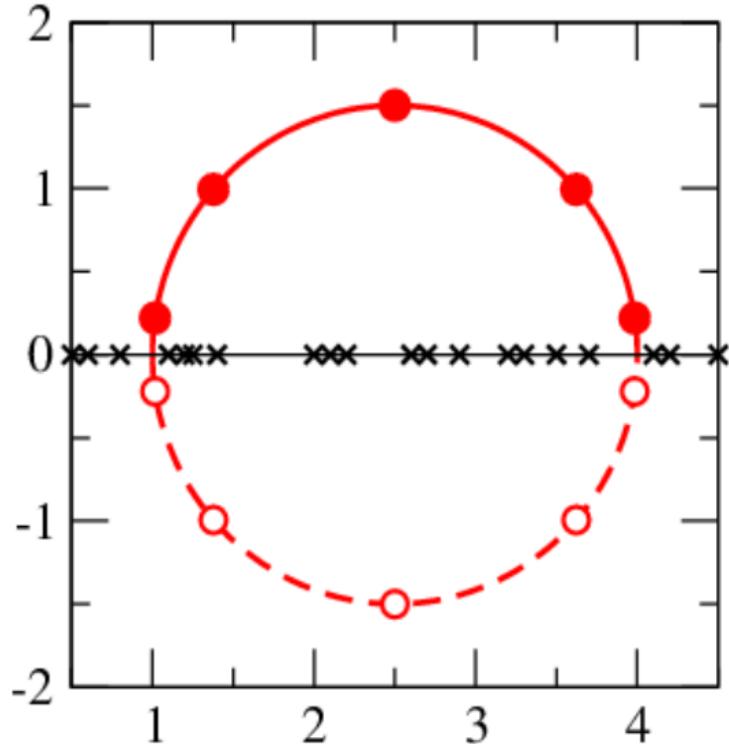
- Interior eigenvalue problems
- Compute eigenvalue by solving linear systems
- Standard/generalized Hermitian/non-Hermitian problems
- Matrix format independent
 - Banded, sparse and dense predefined interfaces
 - Can import a custom linear solver through RCI

www.feast-solver.org

v2.1 included in Intel MKL

FEAST = Subspace Iteration + Spectral Projection

1. Define accelerator



$$Q_{m_0} = \sum_{j=1}^{n_e} \omega_j (z_j B - A)^{-1} B \tilde{X}_{m_0}$$
$$Q_{m_0} = \sum \omega_j Q_{m_0}^{(j)}$$
$$(z_j B - A) Q_{m_0}^{(j)} = B \tilde{X}_{m_0}$$

2. Project matrices

3. Solve reduced Problem

4. Recover approximate eigenpairs

5. Iterate until convergence

$$A_q = Q^H A Q \quad B_q = Q^H B Q$$

$$A_q W_q = B_q W_q \Lambda_q$$

$$\tilde{X} = Q W_q \quad \tilde{\Lambda} = \Lambda_q$$

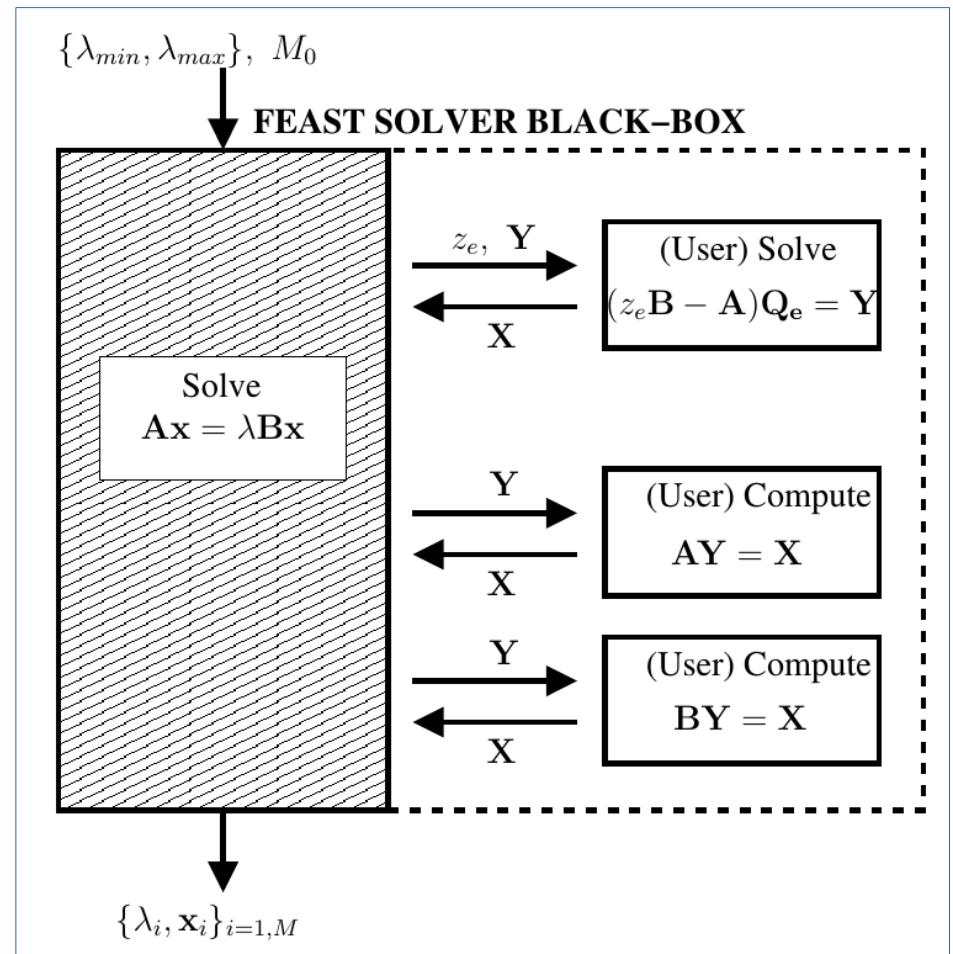
FEAST Pseudo Code and RCI

Algorithm 1 The FEAST Algorithm

```

1: input:  $A, B, \tilde{X}_{m_0}, \{(z_j, \omega_j)\}_{1,\dots,n_e}, \epsilon$ 
2: while ( $\|A\tilde{X}_m - B\tilde{X}_m\Lambda_m\| > \epsilon$ ) do
3:    $Q_{m_0} = 0$ 
4:   for ( $j = 0; j < n_e; j = j + 1$ ) do
5:      $Q_{m_0}^{(j)} \leftarrow (z_j B - A)^{-1} B \tilde{X}_{m_0}$ 
6:      $Q_{m_0} \leftarrow Q_{m_0} + \omega_j Q_{m_0}^{(j)}$ 
7:   end for
8:    $A_q = Q^H A Q \quad B_q = Q^H B Q$ 
9:   Solve  $A_q W_q = B_q W_q \Lambda_q$ 
10:   $\tilde{X}_{m_0} = Q_{m_0} W_q \quad \tilde{\Lambda} = \Lambda_q$ 
11: end while
12: output:  $\tilde{X}_m, \tilde{\Lambda}_m$ 

```



Three Levels of MPI Parallelism

Previous releases of FEAST (v3.0 and prior) offer only 2 levels of MPI parallelism

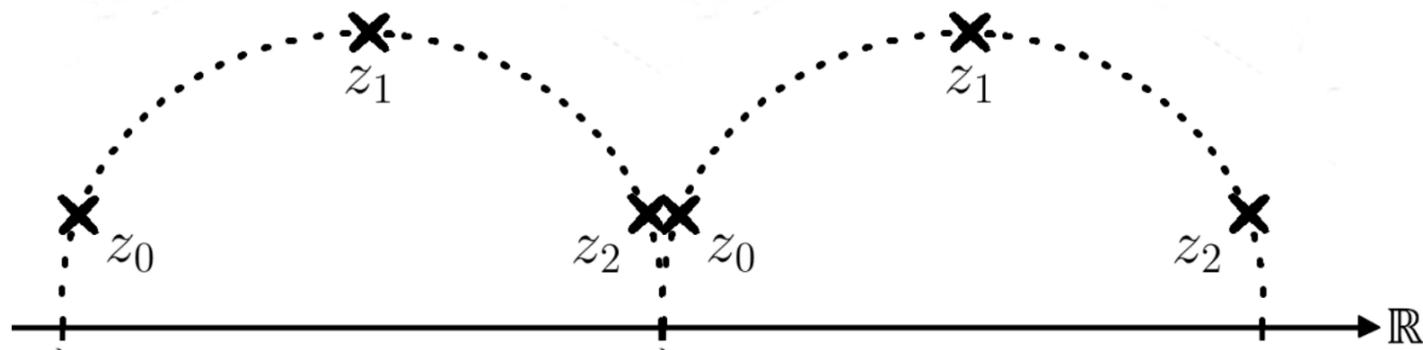
L1: Solving multiple search intervals in parallel (no overlap)

L2: Solving the multiple (independent) linear systems in parallel

L3: Solving each linear system in parallel

First Level of Parallelism

L1: Solving multiple search intervals in parallel (no overlap)



Allows to limit the search subspace size

m_0 should be less than ~ 1000 because of $O(m_0^3)$ scaling of the reduced eigenvalue problem

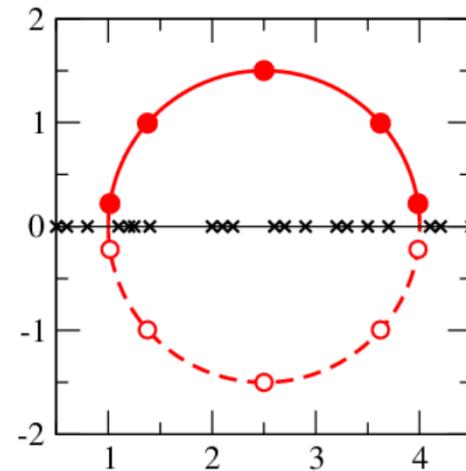
Global orthogonality (between contours)
is largely preserved

Second Level of Parallelism

L2: Solving the multiple (independent) linear systems in parallel

$$Q_{m_0} = \sum_{j=1}^{n_e} \omega_j (z_j B - A)^{-1} B \tilde{X}_{m_0}$$

$$(z_j B - A) Q_{m_0}^{(j)} = B \tilde{X}_{m_0}$$



Advantages:

- Perfect parallelism up to the number of quadrature nodes
- Ideal case #**L2** = n_e → Factorization computed once

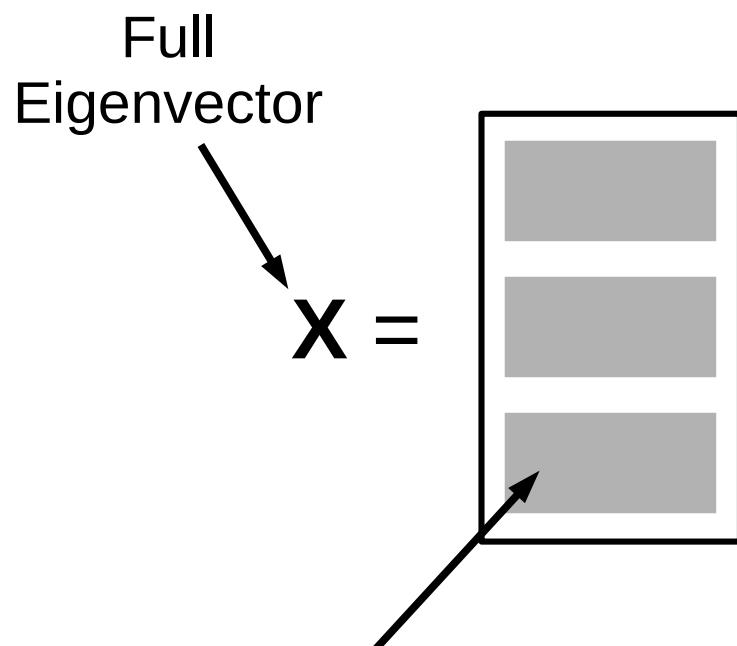
Limitations:

- Memory footprint increased because of extra copies of the matrix and eigenvalues/eigenvectors

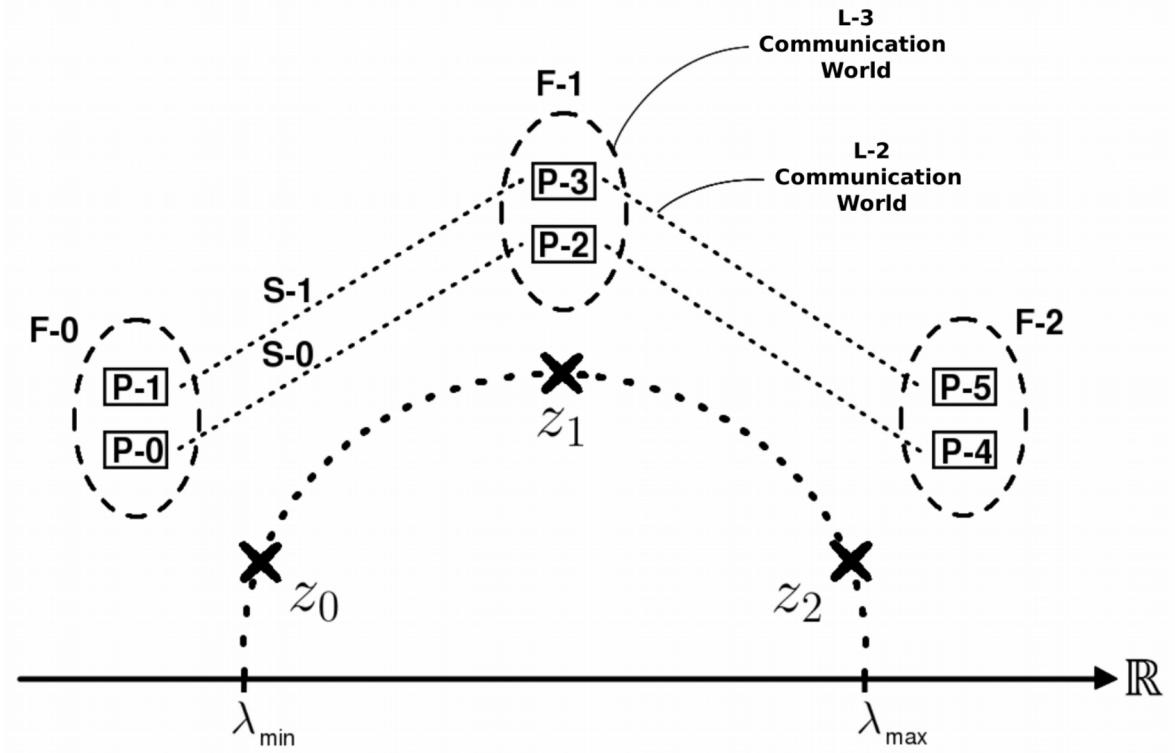
Third Level of Parallelism

L3: Solving each linear system in parallel

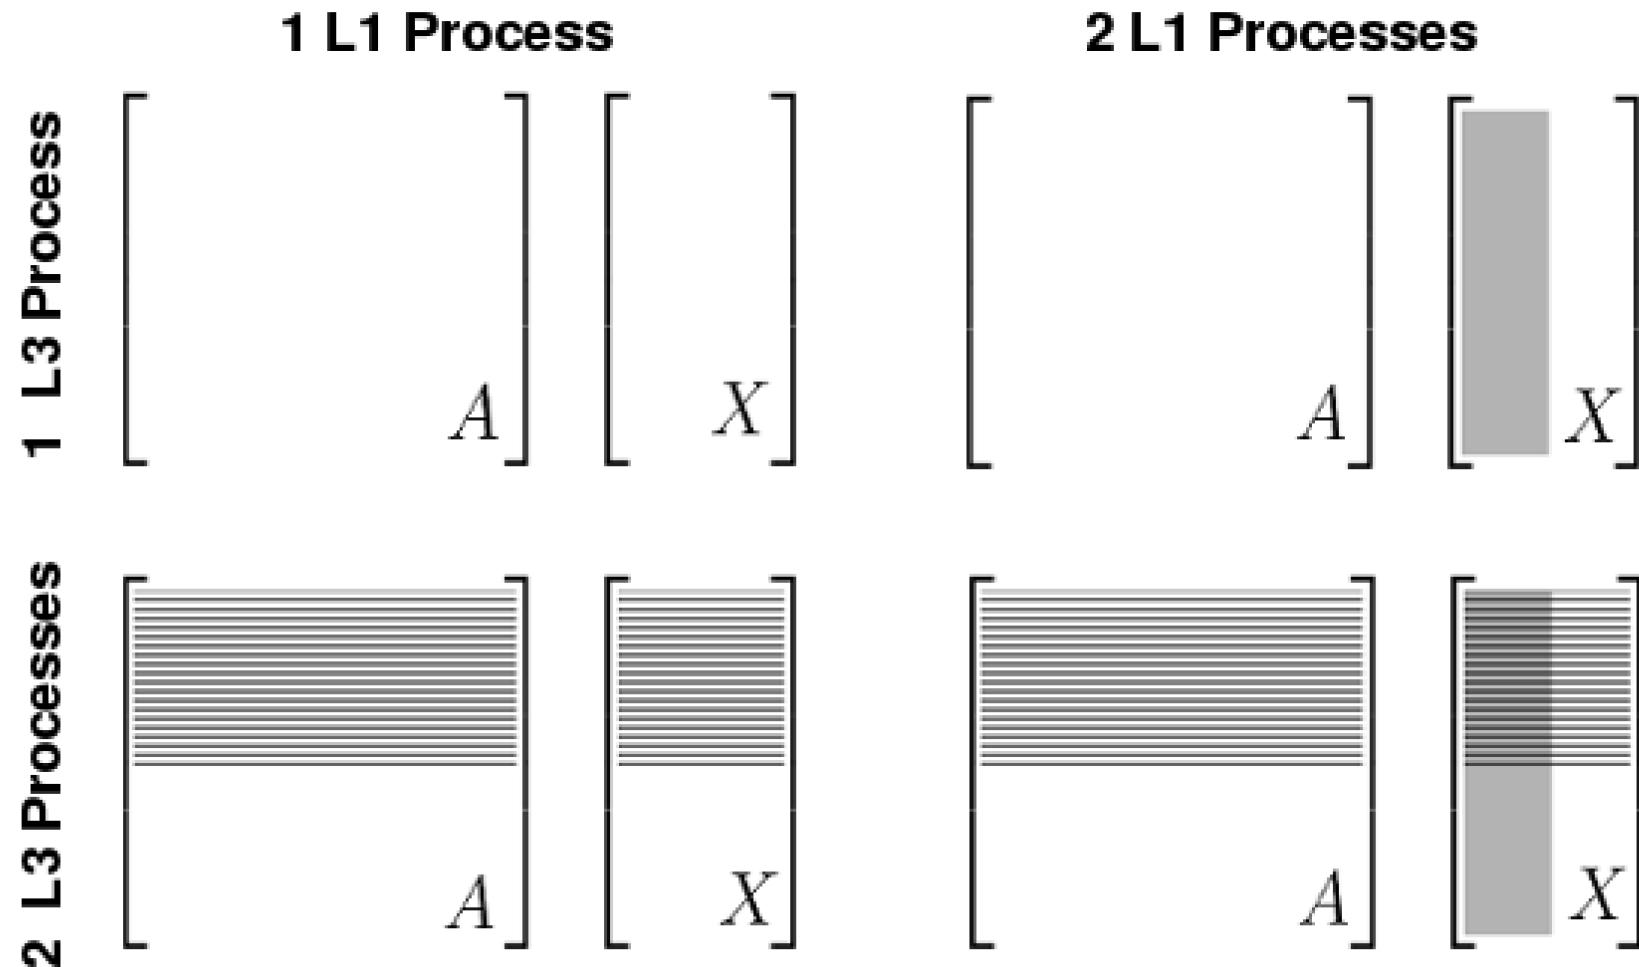
$$(z_j B - A) Q_{m_0}^{(j)} = B \tilde{X}_{m_0}$$



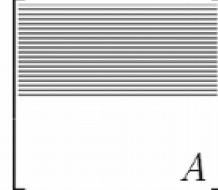
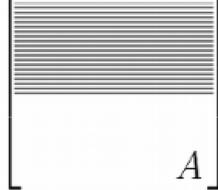
L2 and L3 Communicators must be managed



Reducing the Memory Footprint with L1 and L3



PFEAST with Sparse Direct Solvers

	MKL-Cluster-PARDISO	MUMPS
Matrix	<ul style="list-style-type: none">Distributed 1D by Row 	<ul style="list-style-type: none">Distributed 1D by Row 
RHS/ Solution	<ul style="list-style-type: none">Distributed 1D by Row 	<ul style="list-style-type: none">Fully Constructed on Head Node 

* P.R. Amestoy, I.S. Duff and J.-Y. L'Excellent, Multifrontal parallel distributed symmetric and unsymmetric solvers (1998)

*A. Kalinkin and K. Arturov “Asynchronous approach to memory management in sparse multifrontal methods on multiprocessors” 12

Domain Decomposition Solvers

Reordered Matrix

$$A = \begin{bmatrix} C & E \\ E^T & D \end{bmatrix}$$

Independent blocks

Connections

Linear System to Solve:

$$\begin{bmatrix} C & E \\ 0 & S \end{bmatrix} \begin{bmatrix} X_l \\ X_e \end{bmatrix} = \begin{bmatrix} Y_l \\ Y_e - E^T C^{-1} Y_l \end{bmatrix}$$

Schur Complement
 $(D - E^T C^{-1} E)$

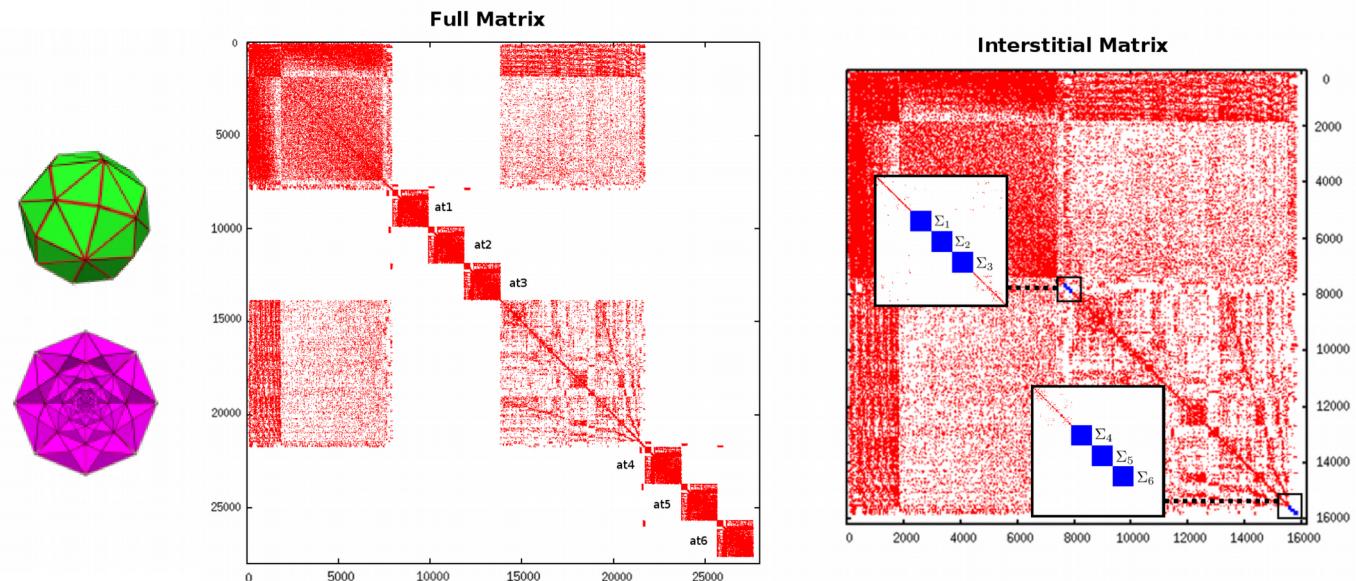
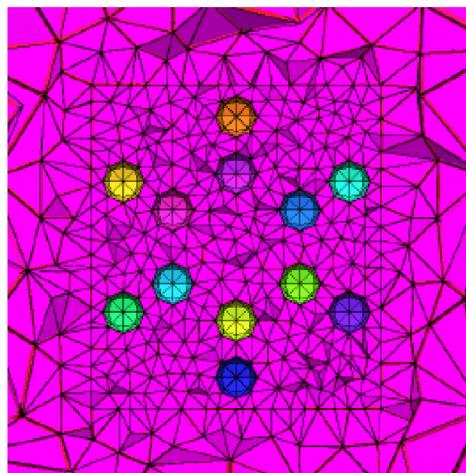
Three step process for Schur

- (i) Solve: $CX_l = Y_l$
- (ii) Solve: $SX_e = Y_e - E^T T$
- (iii) Solve: $CX_l = Y_l - E^T X_e$

*Vasileios Kalantzis “Domain Decomposition Techniques for Contour Integration Eigenvalue Solvers”

Electronic Structure Application

Muffin-tin Domain Decomposition



Use a Schur Complement Approach

- MKL-Cluster-Pardiso (distributed-memory solver) is used to solve the “Interstitial” region
- MKL-Pardiso (shared-memory solver) is used to solve the “Atom” regions

*A. Levin, D. Zang, and E. Polizzi. "FEAST fundamental framework for electronic structure calculations: Reformulation and solution of the muffin-tin problem."

Minnesota Supercomputing Institute

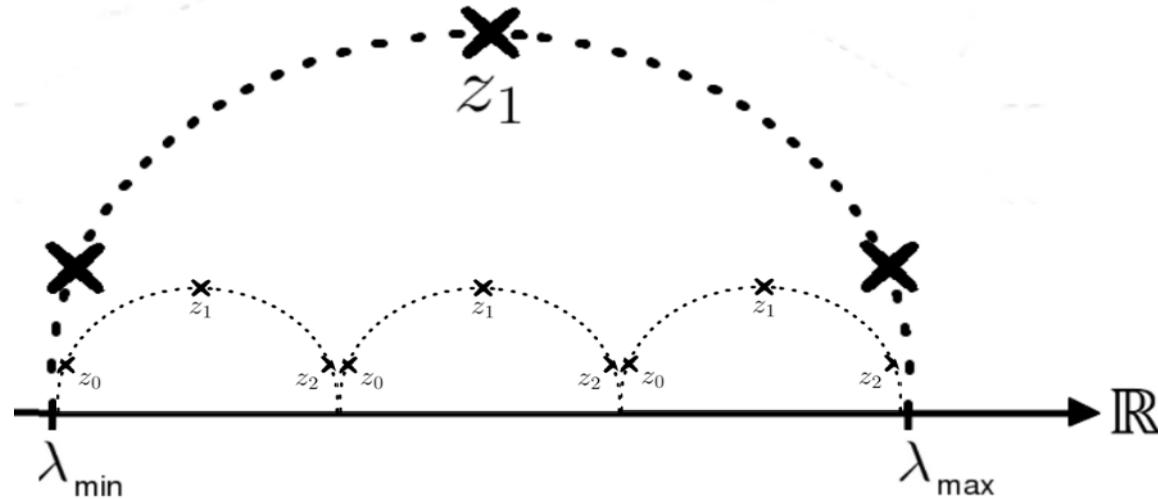
Mesabi Cluster:

- 2 Sockets per Physical Node
 - 2.5GHz Haswell E5-2680v3 12 Core Processor
- 64 GB DRAM
- InfiniBand Interconnect
- Each MPI process defined as a single socket
 - 2 MPI processes per Physical Node
(12 Cores, 12 Threads, 32GB)



www.msi.umn.edu

Results: L1 Strong Scalability



#L1 = 1 and 3

#L2 = 1

#L3 = 13

#MPI: from 13 to 39

#Cores: from 156 to 468

	λ_{min}	λ_{max}	m_0	# Eig	Fact (s)	Solve (s)
C0	-500	-2	600	273	96	170
C1	-500	-100	200	66	94	59
C2	-100	-16	200	92	99	57
C3	-16	-2	200	115	96	56

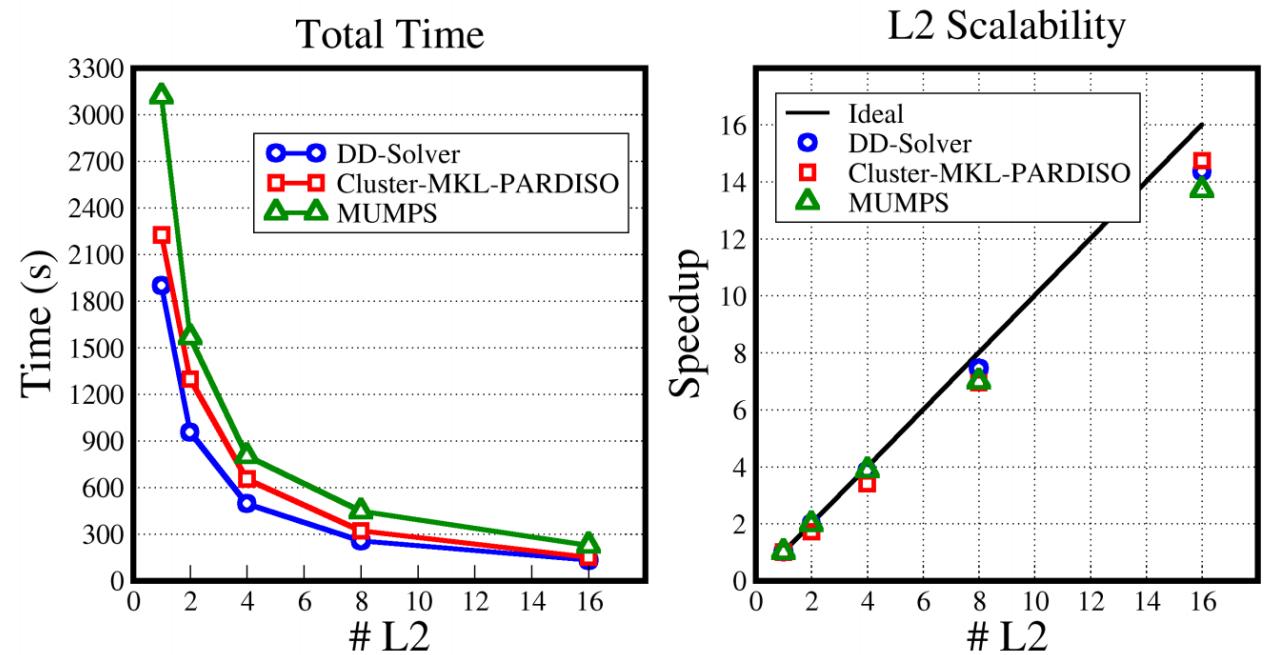
N=302,295; 16 Quadrature Nodes; 1 FEAST Iteration

Results: L2 Strong Scalability

Each independent linear system $(z_j B - A)Q_{m_0}^{(j)} = B\tilde{X}_{m_0}$ can be solved in parallel

#L2 from 1 to 16
#L3 = 3

#MPI from 3 to 48
#Cores from 36 to 576



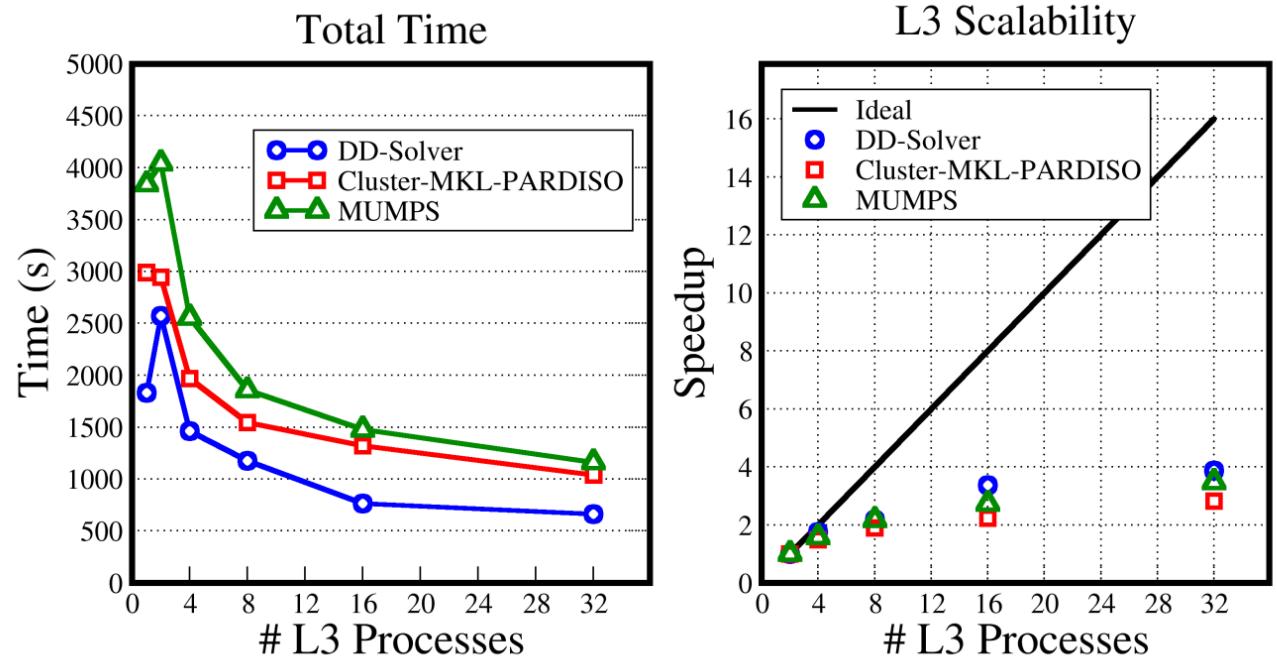
N=302,295; M0=200; 16 Quadrature Point

Results: L3 Strong Scalability

Distributed Memory solver for $(z_j B - A)Q_{m_0}^{(j)} = B\tilde{X}_{m_0}$

#L2 = 1
#L3: from 1 to 32

#MPI: from 1 to 32
#Cores: from 12 to 384

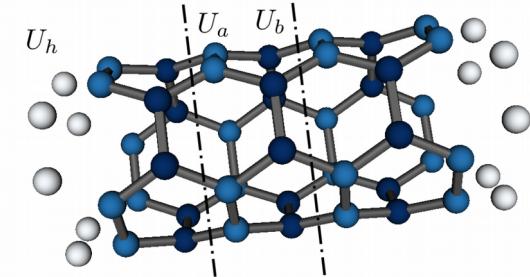
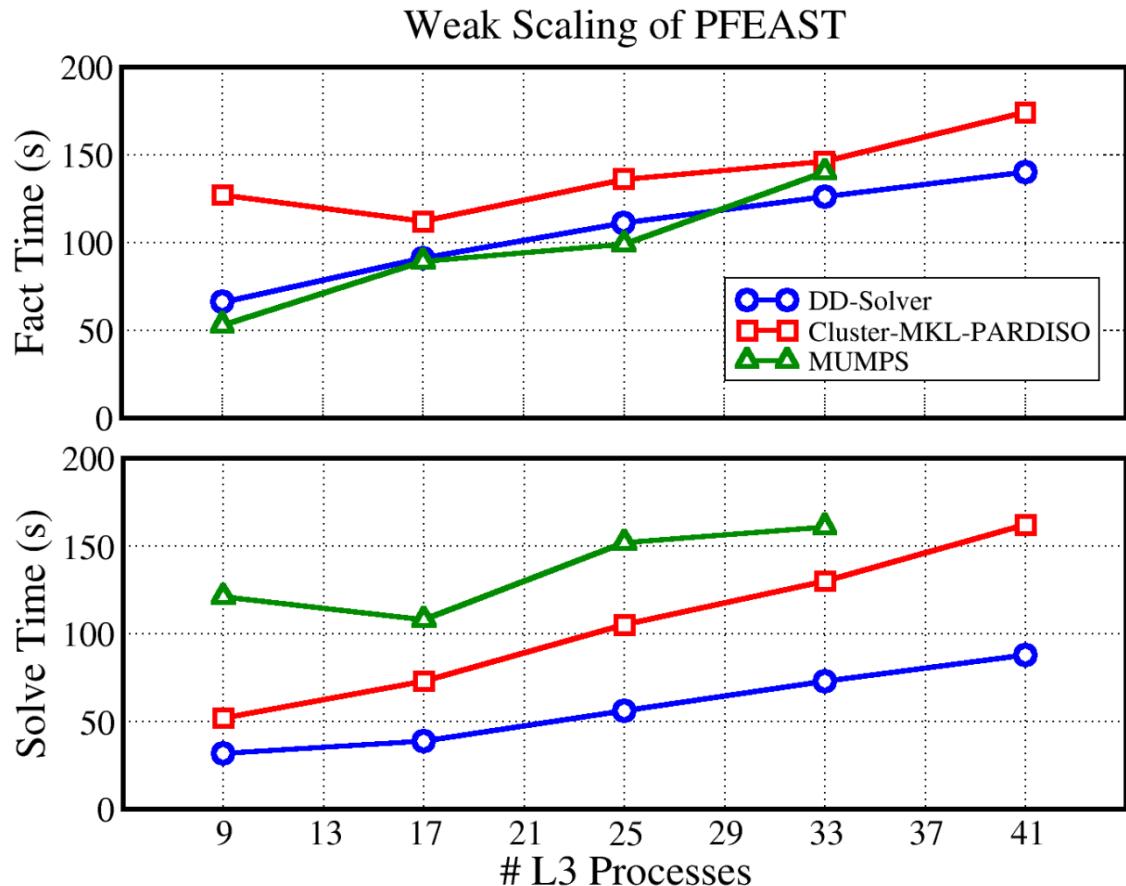


N=302,295; M0=200; 16 Quadrature Point

Results: Weak Scalability

# MPI	9	17	25	33	41
# Atoms	54	102	150	198	246
System Size	211,110	392,650	575,760	757,934	942,157

M0=200, 16 Contour Points (no L2); 1 FEAST Iteration



Example 246 Atom CNT

Need 750 eigenvectors
#L1=5 → 5 slices M0=200

#L2=16 (Optimal Case)
#L3=41

$$\begin{aligned}
 & (\#L1) \times (\#L2) \times (\#L3) \\
 & = 3,280 \text{ MPI} \\
 & = 39,360 \text{ Cores}
 \end{aligned}$$

Conclusion

FEAST v4.0 (PFEAST) Coming Soon!

- 3 Levels of MPI Parallelism
- Reduce Memory with **L1** and **L3**
- 1D Eigenvector Distribution
- RCI for Custom Solver
- Predefined Interfaces for MUMPS and MKL-Cluster-Pardiso

- Scale to Larger Molecules (up to 1000 atoms) with Electronic Structure Application

